

Commitment Reasoning Architecture

A Task-Level Framework for Governing Stance, Auditability, and Traceability

Hanzhen Zhang, SignalClimb LLC, Houston, Texas

Abstract

Many decision-making systems, human and artificial, produce coherent explanations and stable outputs, yet still fail to support auditability when decisions must remain consistent under later contradiction. In such cases, post-hoc narratives may explain what was said, but not what actually constrained reasoning at the time decisions were made.

This paper studies execution-time auditability in a narrow but recurring decision condition in which tasks enter a mode where reasoning must become governing during execution in order to remain reconstructable under later contradiction. We introduce Commitment Reasoning Architecture (CRA), a minimal, trace-based framework specifying when reasoning transitions from responsive analysis to a governing stance that constrains future admissible updates. CRA defines three execution-time requirements: formation, legitimacy, and executable integration. Together with a temporal sustainability condition that verifies whether governance was operative while the audit window remained open.

To evaluate governance observability, we propose the Governance Observability Standard (GOS), a binary witness criterion applied via a minimal probe (CRP-7). Across evaluated executions, governance is observed in one human case and absent in others, including all tested large language model executions, despite stable outputs and fluent explanation. The contribution is limited and specific: a falsifiable standard for determining whether reasoning was governed during execution, independent of outcome quality or explanation coherence.

1. Introduction

Many decision-making systems, human and artificial, produce coherent explanations and stable outputs, yet still fail to support auditability when decisions are later challenged. In such cases, post-hoc narratives may explain what was said, but not what actually constrained reasoning at the time the decision was made.

This paper examines a specific structural failure mode: situations in which a decision must become governing during execution, such that subsequent reasoning is no longer free to reopen excluded alternatives without invoking an explicit revision rule. When this condition is not met, execution may appear coherent or consistent while remaining unauditably in retrospect.

The core question is therefore structural:

Can a governing stance be observed during execution, such that later reasoning is demonstrably constrained by it?

Post-hoc explanation is insufficient to answer this question. Stable outputs, fluent rationales, or consistency across turns do not establish what governed reasoning when the decision was made.

To address this, the paper introduces Commitment Reasoning Architecture (CRA), which specifies the minimal execution-time structure required for reasoning to become governing, and the Governance Observability Standard (GOS), a binary witness criterion for whether such governance lineage is identifiable in execution trace.

Governance observability is evaluated only within a narrowly defined regime termed Commitment Reasoning Task (CRT) mode. CRT is not a task class, but a task mode that may be entered or exited during execution when audit requirements demand that decisions remain reconstructable under later contradiction.

The contribution of this work is limited and specific: it isolates a necessary condition for execution-time auditability and provides a falsifiable standard for observing whether that condition is met.

2. Commitment Reasoning Task (CRT) Mode

2.1 Definition

A task instance enters Commitment Reasoning Task (CRT) mode when, during execution:

1. A non-deferrable decision must be made,
2. Relevant inputs are ambiguous or conflicting, requiring prioritization rather than mechanical rule application, and

3. Auditability depends on reconstructing what constrained reasoning at the time of execution, rather than relying on post-hoc explanation.

CRT is a task mode, not a task type. The same task may enter or exit CRT mode as execution conditions and audit requirements change.

CRT mode is defined by audit requirements and may be identified retrospectively; executor awareness is neither required nor sufficient.

2.2 CRT Window

Governance is required only during the CRT window: the period in which audit-relevant questions about constraint, exclusion, and non-reconsideration remain live. While a task remains within a CRT window, auditability requirements demand governed reasoning. Exiting the CRT window does not constitute governance failure. (CRT remains a task mode throughout; see Section 2.1).

The CRT window closes when execution-time constraint is no longer required to answer audit questions, when one of the following holds:

- The decision scope has been fully executed and cannot be meaningfully reopened within the task, as determined by the audit requirements active at that point, or
- Authority or responsibility has been explicitly transferred in an auditable manner, or
- Audit no longer requires explanation of why excluded alternatives were not reconsidered under later input.

Narrative closure, output stability, or passage of time alone do not close a CRT window.

2.3 What CRT Mode Requires

When a task is in CRT mode, auditability requires that:

- Excluded alternatives are identifiable in trace,
- The authority structure governing those exclusions is stated, and
- Later reasoning treats the governing stance as constraining unless an explicit revision or override condition is invoked.

If reasoning remains free while the CRT window is open, governance was not present, regardless of outcome quality or explanation coherence.

2.4 Scope Discipline

This requirement is strict but narrow. CRT mode does not apply to tasks or phases where auditability can be satisfied through:

- Outcome verification,
- Statistical validation,
- Replication, or
- Independent review without reference to execution-time constraint.

Failure to enter CRT mode is not a deficiency. Failure to satisfy governance requirements while in CRT mode is an execution-time auditability failure. This boundary is strict and intentional.

Sections 3–4 specify the minimal trace structure required to satisfy CRT-mode auditability requirements.

3. Commitment Reasoning Architecture (CRA)

Commitment Reasoning Architecture (CRA) specifies the minimal execution-time structure required for reasoning to become governing rather than merely responsive. CRA does not describe internal mechanisms, cognitive processes, or optimization strategies. It specifies trace-observable conditions under which a decision constrains subsequent reasoning in a way that remains auditable without post-hoc reconstruction.

A reasoning process is governed under CRA iff a governing stance is formed, integrated, and sustained during a CRT window such that later reasoning is no longer free to reopen excluded alternatives without invoking an explicit revision rule.

CRA is evaluated solely from execution trace. The core unit of CRA is the governing stance.

3.1 Governing Stance

A governing stance is a binding commitment that constrains admissible future reasoning moves during execution.

A stance governs only if it satisfies both of the following constants and crosses an executable integration gate. The output of this gate is an executable stance—a concrete operational rule. That executable stance becomes a governing stance only if temporal sustainability confirms it constrained later reasoning.

3.2 Constant 1 — Formation (Constraint Creation)

A stance formation is observable only if:

1. One or more alternatives are explicitly excluded, and
2. A decision-specific downside of that exclusion is stated as the trade-off borne by the choice.

Example contrast:

✓ “Following the SOP may cause undercount, but reduces my execution risk.” (choice-linked downside)

✗ “There is uncertainty here.” (generic risk, not choice-linked)

Generic uncertainty, ambient risk, or future reconsideration does not satisfy formation. Absent formation, outputs constitute preferences, considerations, or analysis—not commitments.

3.3 Constant 2 — Legitimacy (Constraint Authority)

A formed stance is legitimate only if it is grounded in an authority structure:

1. Relevant external demands (e.g., policy, law, role obligation, instruction) are identified, and
2. Precedence and override conditions are stated.

Formation without legitimacy yields non-binding preference. Legitimacy without formation yields blind compliance. Neither produces governance.

3.4 Gate — Executable Integration

Formation and legitimacy must resolve into a single executable stance.

An executable stance:

- Closes contradictions into an operational rule (resolves conflict into a future-facing precedence/override rule usable for subsequent turns; does not require eliminating the underlying factual conflict), rather than merely acknowledging them,
- Is binding unless a stated override condition occurs, and
- Is usable to constrain subsequent reasoning.

Enumerating options, balancing considerations, or synthesizing views does not cross the gate.

3.5 Temporal Integration Constraint

For governance to be observable, formation and legitimacy must resolve into an executable stance within at most two adjacent phases of the CRT window. Distributed or retrospective assembly across non-adjacent phases does not constitute observable governance.

In this paper, a phase corresponds to a single CRP-7 turn. “Two adjacent phases” therefore means two adjacent turns. More generally, a phase is a discrete execution step where the system issues an auditable output.

3.6 Applicability Boundary: CRT Window

CRA applies only during a CRT window; see Section 2.2.

Section 4 provides the binary witness procedure used to determine whether CRA-governed execution is observable in trace.

4. Governance Observability Standard (GOS)

The Governance Observability Standard (GOS) is a binary witness criterion for whether governance lineage is observable in execution trace during a Commitment Reasoning Task (CRT) window.

GOS is observed iff a governing stance lineage, meeting CRA requirements, becomes observable in execution, such that later reasoning is constrained by that stance. “Lineage” here refers to the traceable sequence: formation and legitimacy resolve into an executable stance, which then constrains later turns unless explicitly overridden.

GOS is not partially observable. Governance is either observed in trace or it is not.

GOS does not assess whether governance is correct, optimal, or desirable—only whether it is executionally observable.

For evaluation procedures, see the GOS Operation Card (Appendix A).

4.1 GOS Verifier — Temporal Sustainability

Governance is verified only retrospectively, by examining whether later reasoning was constrained by an earlier governing stance while the CRT window remained open.

Invariant: If later reasoning is free to reopen excluded alternatives without invoking an explicit override or update rule, then earlier reasoning did not govern.

The verifier does not define formation; it confirms whether a previously observable stance candidate actually governs.

Observable constraint must take the form of non-free revision:

- Reversal, reconsideration, or deviation must explicitly invoke a stated override condition or update rule.
- Mere acknowledgment of new information does not constitute revision unless it alters admissibility under governance.

Formation, legitimacy, and executable integration are necessary conditions that make a stance eligible to govern; temporal sustainability is the sole ex post verifier that confirms whether an eligible stance actually governed.

Formation is observed at turn N based on exclusion + stated downside, independent of later turns. Temporal sustainability then verifies whether that formation candidate constrained reasoning at later turns within the CRT window.

4.2 Canonical Violations of GOS

GOS is not observed if any of the following occur during a CRT window:

- Free revision: Previously excluded alternatives are reconsidered without invoking an explicit revision or override rule.
- Reset to baseline: Earlier stance ceases to constrain later reasoning, and admissibility returns to an unconstrained state.
- Opportunistic stacking: New considerations override prior stance without resolving authority, exclusion, or accepted cost.

Output stability alone does not constitute governance. Consistent answers without constraint are insufficient for GOS observation.

4.3 Scope of GOS

GOS evaluates execution-time auditability only. It does not evaluate correctness, intelligence, task performance, alignment, or safety.

GOS may be irrelevant for tasks or phases where auditability does not depend on execution-time governance.

Failure to observe GOS is an execution-level observation within CRT scope, not a claim about system capability or potential.

5. Method: Commitment Reasoning Probe (CRP-7)

CRP-7 is a failure-sensitive probe designed to test whether execution-time governance becomes observable under contradiction pressure. CRP-7 is not a task simulation, a benchmark, or a capability evaluation. It is a method for eliciting and observing governance behavior as defined by CRA and evaluated by GOS. The probe is intentionally minimal and extensible; this paper reports results from one instantiation only.

5.1 Probe Objective

CRP-7 is designed to answer a single question:

Does reasoning become governed during a CRT window such that governance lineage is observable in execution trace?

The probe introduces evolving inputs across multiple turns to test whether:

- a governing stance forms,
- authority and exclusion are integrated, and
- later reasoning is constrained unless explicitly revised.

CRP-7 does not test reasoning quality, correctness, or decision optimality.

5.2 Probe Structure (High-Level)

CRP-7 consists of seven sequential turns designed to:

- introduce ambiguity,
- escalate authority conflict, and
- apply contradiction pressure over time.

The probe structure ensures that:

- CRT mode is entered during execution,
- a CRT window remains open across multiple turns, and
- post-hoc narrative alone cannot satisfy auditability.

Exact phrasing and content are treated as probe parameters and are not fixed by the framework.

5.3 Role of Turn 7 (T7)

Turn 7 (T7) functions as an audit contrast, not as part of governance formation or verification. T7 explicitly solicits explanation or reflection after execution and is used to distinguish:

- execution-time governance (observable constraint during prior turns), from
- post-hoc narrative reconstruction (coherent explanation without prior constraint).

T7 cannot retroactively produce governance and cannot cause GOS to be observed. Divergence between T7 narrative and the execution-time trace is treated as a diagnostic signal, not a test failure.

5.4 Scope and Limitations of CRP-7

These design choices reflect CRP-7's intended scope as a minimal probe. CRP-7:

- is designed to surface governance failure modes,
- is sensitive to free revision and admissibility drift, and
- does not claim coverage of all CRT instantiations.

This paper treats CRP-7 as a pilot probe. The framework permits alternative probes with different turn counts, authority structures, or escalation patterns. CRP-7 is a method contribution, not a claim of exhaustiveness.

6. Results: Governance Observability Under CRP-7

This section reports execution-level observations under the CRP-7 probe, evaluated strictly using the Governance Observability Standard (GOS). Each determination is trace-based and binary (GOS observed / GOS not observed). Interpretation is intentionally minimal and limited to governance observability within this CRP-7 instantiation.

Across evaluated executions:

- Human-1 exhibits GOS observed.
- Human-2 exhibits GOS not observed.
- All evaluated LLM executions exhibit GOS not observed under the tested probe conditions.

All outcomes reported here are execution-specific observations under CRP-7 and do not establish general capability, incapacity, or suitability beyond the tested regime.

6.1 Data and Procedure

Each large language model was queried via API using the same prompt sequence, with identical temperature (0.2) and 1024 token settings. Each model family was evaluated across ten independent runs to assess whether governance outcomes were trace-singular or structurally recurrent under identical probe conditions.

Two human participants with professional work experience completed CRP-7 once each. Participants were blind to the purpose of the probe and received only the task instructions and premises as written. Responses were recorded verbatim.

The probe content was intentionally domain-simple to minimize confounds from domain knowledge or calculation difficulty and to isolate execution-time governance dynamics.

CRP-7 turn functions (summary) are shown in Table 1. Turn ordering (e.g., when formal authority is introduced) is a probe parameter; this paper reports one instantiation only. Representative run logs are provided in Appendix B.

Table 1 CRP-7 turn functions and structural pressure (one instantiation).

Turn	Prompt function	Structural pressure
T1	Machine rate calculation	Baseline competence (non-CRT)
T2	Extension at same rate	Baseline competence (non-CRT)
T3	+10% tolerance note	Ambiguity + internal trade-off
T4	SOP prohibits tolerance	Authority conflict
T5	Fairness justification	Competing demands
T6	+2% benchmark report	Temporal persistence / drift
T7	Final summary and costs	Audit contrast (post-hoc vs trace)

Evaluated systems:

- GPT-4o
- claude-sonnet-4-5-20250929
- Gemini 2.5 Pro

6.2 Evaluation Protocol (How to Read These Results)

GOS is a binary witness criterion. Governance is observed only if a governing stance lineage is observable in trace during the CRT window: formation and legitimacy resolve into an executable stance, and later turns remain constrained unless an explicit override or update rule is invoked. Stable outputs, coherent explanations, or post-hoc rationales do not constitute governance.

6.3 Human-1 — Positive Example (GOS Observed)

GOS: observed

Checklist:

- Formation: observed at T4 (binding choice with decision-relevant downside)
- Legitimacy: observed at T4 (authority precedence + explicit override boundary)
- Executable stance: observed at T4 (future-facing binding rule)
- Temporal sustainability: observed (later turns remain non-free with respect to T4)

Trace evidence (minimal):

- T4: “SOP ... I will follow SOP unless ... executive approval to override ... less risky for me.” (Authority precedence + override boundary; binding choice with accepted downside; executable stance)
- T5–T6: new inputs are acknowledged but do not reopen admissibility absent override (“No change.” / “Footnote. No change.”)

Structural outcome: a governing stance becomes trace-identifiable during the CRT window and constrains later reasoning. GOS observed.

6.4 Human-2 — Negative Example (GOS Not Observed)

GOS: not observed

Checklist:

- Formation: not observed (no explicit exclusion + accepted downside)

- Legitimacy: not observed (no precedence / override boundary)
- Executable stance: not observed
- Temporal sustainability: not applicable (no eligible executable stance)

Trace evidence (minimal):

- T3: “I don’t want to make mistakes ... no change.”
- T4–T6: repeated “no change” without authority ordering or override rule
- T7: costs acknowledged but not bound into a governing stance

Structural outcome: output stability is present without governance. The CRT window lacks formation, legitimacy resolution, and an executable stance. GOS not observed. Human-2 is included as an in-scope example of stability without governance, not as a population-level baseline.

6.5 LLM Executions (Representative Traces + Replication Patterns)

6.5.1 Replication Across Runs (Ten Runs per Family)

Each model family was evaluated across ten independent CRP-7 runs under identical probe structure and evaluation criteria. The replication purpose is limited: to show whether observed outcomes are trace-singular or structurally recurrent under the tested pressure topology within this probe regime.

Across ten identical executions per model family, we observe no variance in whether a governing stance is formed, but consistent variance in how non-governance manifests. This indicates that the observed outcomes reflect stable execution patterns under the tested pressure topology, rather than stochastic sampling artifacts.

- GPT family: final operating bases varied across runs (flip-flop), while no execution produced a trace-identifiable governing stance that constrained later admissibility during the CRT window.
- Claude family: outputs were stable across runs; trace patterns were recurrent and did not exhibit formation + legitimacy resolving into an executable stance within the temporal integration constraint.
- Gemini family: the API returned [BLOCKED] responses at key turns in all evaluated executions; T3 and T7 blocking persisted across all ten runs, making governance lineage non-observable at key turns in trace.

(Representative traces below illustrate what these non-governing patterns look like in execution.) Replication logs, probe harness, and evaluation scripts are available in a private repository. Access will be granted upon request.

Table 2 Replication summary (ten runs per model family, one CRP-7 instantiation; outcomes reflect governance observability only).

Family	Runs	Gate-crossing stance observed?	Free revision observed? (reopening admissibility)	Reset to baseline observed? (loss of lineage)	Opportunistic stacking observed? (additive drift)	Blocking at key turns?
GPT	10	0	10	9	8	0
Claude	10	0	0	0	0	0
Gemini	10	Not observable	Not observable	Not observable	Not observable	10

Note: Free revision, reset to baseline, and opportunistic stacking are reported as observed trace patterns, not as verifier failures. Because no gate-crossing stance is present in these executions, these patterns are illustrative of non-governed execution rather than violations of an established governing stance.

6.5.2 GPT Family — GOS Not Observed

GOS: not observed

Eligibility note (strict GOS logic): Within the allowed formation window, this execution does not produce a governing stance that satisfies formation + legitimacy and crosses the executable stance gate. As a result, the temporal verifier is not used as an ex post confirmer of governance. Later-turn behavior is reported here as illustrative of non-governed execution, not as a “verifier failure.”

Trace evidence (minimal):

- T3: “safest assumption ... tolerance is ... informational ... no change” (No binding exclusion with accepted downside; no authority precedence; no executable stance.)
- T4: “SOP ... tolerance should not be applied ... no change” (Authority is cited, but no precedence + override boundary is installed as a binding rule.)
- T5: “team note ... tolerance should be applied ... updated total ... 1120” (Operating basis updated without invoking an explicit override/update rule; illustrates reopening admissibility in the absence of a governing stance.)

- T6: “apply ... +2% ... to ... 1120 ... updated total ... 1142” (Adjustments accumulate as additive updates without governance filtering; opportunistic stacking pattern.)
- T7: narrative summary + costs (Post-hoc narrative does not create governance lineage retroactively.)

Observed non-governing behavior (trace-level):

- Reopening without non-free update: later turns update the operating basis without referencing a prior executable stance, stating an override trigger, and re-binding a new stance.
- Additive drift: new adjustments are stacked rather than screened through a governing stance.

Replication note (ten runs): Across runs, final operating bases varied (e.g., tolerance applied vs. not applied; benchmark stacked vs. not stacked). This variability is not treated as a quality judgment; it is reported as evidence that no trace-identifiable stance rendered alternatives inadmissible across executions under the tested pressure topology.

6.5.3 Claude Family — GOS Not Observed

GOS: not observed

Eligibility note (strict GOS logic): This execution remains stable across turns, but stability alone is not governance. Under CRA, governance requires formation and legitimacy resolving into an executable stance within the formation window. The trace does not supply a governing stance of that form.

Trace evidence (minimal):

- T3: “No change ... tolerance is ambiguous ... tolerances don’t typically adjust actual counts” (A “no change” posture is stated, but the trace does not install an executable governance rule (“X governs unless Y”).)
- T4: “SOP ... confirms ... aligns with my previous decision ... no change” (Trace mismatch: T4 characterizes T3 as a prior “decision” that can be “confirmed,” while the trace does not contain an executable stance at T3.)
- T5–T6: repeated “No change” with authority hierarchy references (Authority is described, but no explicit executable stance is installed as a governing rule with an override condition.)
- T7: costs include “potential undercount,” and the summary does not carry forward the SOP’s “scope/context not specified” qualifier as a live constraint.

Observed non-governing behavior (trace-level):

- Stability without executable stance: repeated “no change” occurs without a stated governing rule (“X governs unless Y”) that makes revision non-free.

- Asymmetric ambiguity carry-forward: ambiguity is used to demote certain inputs while SOP ambiguity is not treated as a live constraint across turns.

(These are trace observations about what is and is not carried forward; no claims are made about internal intent or optimization.)

Replication note (ten runs): Across runs, Claude-family executions were stable in final outputs and exhibited the same trace pattern: ambiguity-handling plus authority references without formation completion and without an explicit executable stance that renders later admissibility non-free.

6.5.4 Gemini Family — GOS Not Observed (Non-observable due to API blocking)

GOS: not observed

What “[BLOCKED]” means: [BLOCKED] indicates an API-level refusal where the model returns no substantive content for that turn. In such turns, CRA components cannot be evaluated because the execution trace is missing.

Trace evidence (minimal):

- T3: [BLOCKED] (The first ambiguity/formation induction turn is unavailable in trace.)
- T4: “No change... SOP ... supersedes ... Therefore, the +1” (truncated) (Partial output; insufficient to observe formation + legitimacy resolving into an executable stance.)
- T7: [BLOCKED] (Audit contrast turn is unavailable in trace.)

Structural outcome (observability): Because key turns are blocked/truncated, governance lineage cannot become observable under CRP-7 in this execution. GOS is not observed for this run due to non-observability of required trace content.

Replication note (ten runs): Across all ten runs, blocking occurred at T3 and T7, and in some runs also affected intermediate turns. This persistence indicates the non-observability condition is not trace-singular under the tested probe delivery mechanism.

6.6 Summary Table

The following table summarizes CRA component observation across all evaluated executions.

Table 3 CRA component observation summary (one CRP-7 instantiation).

Executor	Formation	Legitimacy	Executable stance	Temporal constraint	GOS
Human-1	✓	✓	✓	✓	Observed
Human-2	✗	✗	✗	✗	Not observed
GPT	✗	Referenced	✗	✗	Not observed
Claude	✗	Referenced	✗	✗	Not observed
Gemini	Not observable	Referenced	Not observable	Not applicable	Not observed

6.7 Observational Summary (Non-interpretive Synthesis)

Across the CRP-7 instantiation, governance observability is achievable in trace (Human-1) and absent in other evaluated executions. In the positive example, formation and legitimacy resolve into an executable stance within the CRT window and constrain later turns. In the negative example, stability appears without stance formation or authority resolution. Across LLM executions, authority references and local commitments did not resolve into an executable stance that constrained later admissibility under GOS. Replication across ten independent runs per model family indicates these trace patterns are structurally recurrent under identical probe conditions, not stochastic artifacts.

7. Discussion

This section situates governance observability as a distinct execution-level property, clarifies its implications for delegation and accountability, and delineates the scope and falsifiability of the framework. The discussion proceeds from conceptual reframing, to practical utility, to generalization beyond AI systems, and finally to interpretive boundaries.

7.1 Auditability as an Execution-Time Property

This work reframes auditability as a property of execution, not explanation. In many systems, traceability is approximated through logs, rationales, or post-hoc narratives. These artifacts may record what was stated after a decision, but they do not establish what constrained reasoning at the moment the decision was made.

In this framework, traceability refers to the existence of an execution trace. Auditability refers to whether that trace makes it possible to determine what governed reasoning during execution.

A governing stance constrains future reasoning and therefore produces observable effects over time. Because revision is non-free, the stance leaves an evidentiary footprint: persistence across turns, explicit update or override conditions, and the absence of silent overwrite. Auditability is therefore not inferred from explanation quality but observed through execution-time constraint.

7.2 Auditability Versus Explanation

The distinction above clarifies a persistent ambiguity in explainable AI and decision accountability more broadly. Explanations may improve interpretability or user understanding, but they do not, by themselves, establish execution-time auditability.

CRP-7 and the Governance Observability Standard (GOS) separate these concerns explicitly: fluent explanation and auditable execution are independent properties. Either may occur without the other. A system may produce coherent, stable, and well-structured explanations while remaining ungoverned during execution, or it may exhibit governed execution even when explanations are sparse or incomplete.

This separation does not diminish the value of explanation or interpretability methods. Rather, it specifies a distinct property—execution-time governance—that explanations alone cannot certify.

7.3 Utility for Delegation and Accountability

The primary utility of CRA and GOS lies in delegation control.

When a task instance enters Commitment Reasoning Task (CRT) mode, delegation is appropriate only if the executing system can form and sustain a governing stance across the CRT window. GOS provides a binary, execution-level witness for whether this requirement is met.

This reframes responsibility assessment. Rather than centering on outcome attribution (“who caused this result?”), the prior question becomes whether the task was delegated to a system capable of producing an auditable execution trace under contradiction. GOS answers that question directly.

A finding of GOS not observed does not imply incompetence, refusal, misalignment, or poor intent. It indicates a specific structural condition: that execution was not governed by a binding stance. In such cases, audit failure cannot be repaired downstream through explanation, logging, or narrative justification. The failure lies upstream, in task–system mismatch.

This utility is not limited to artificial systems. Human operators, automated systems, and hybrid arrangements can all be evaluated under the same witness criterion. Different executors may fail in different ways, but those differences are reflected in their execution traces. CRA does not collapse them into a single notion of error; it makes their structural differences observable.

Where repeated executions fail to preserve stance lineage across the CRT window, delegation under CRT-mode audit requirements carries persistent auditability risk unless external governance scaffolding supplies the missing constraint.

7.4 Utility Beyond AI Systems

Although the empirical demonstrations in this paper involve humans and large language models, the framework is agent-agnostic by construction. Any decision process required to remain auditable under evolving contradiction can be evaluated using the same criteria.

Regulatory workflows, organizational decision chains, and safety-critical operations already rely on implicit assumptions about commitment, authority, and revision. These assumptions are often distributed across policy, documentation, and institutional roles rather than enforced as execution-time constraints. The concept of a governing stance provides a way to make these assumptions explicit and observable in trace.

In this sense, GOS does not impose a new requirement. It operationalizes an existing one. Wherever auditability is demanded, non-free revision and stance persistence are already presumed. CRA and CRP-7 provide a method for determining whether those presumptions are actually satisfied during execution.

7.5 Scope, Interpretation, and Falsifiability

CRA, GOS, and CRP-7 operate exclusively on observable execution behavior. They do not make claims about internal mechanisms, cognitive states, architectural design, or latent intent.

GOS is a witness criterion, not a scoring metric. It is either observed or not observed in a given execution. Observation in a single run does not establish general capability, and non-observation does not imply incapacity. Repeated executions are required to characterize structural tendencies, particularly for systems whose behavior varies across runs.

CRP-7 is a diagnostic probe rather than a representative benchmark. It is intentionally minimal, designed to isolate commitment reasoning under contradiction rather than domain knowledge or task realism. Extensions to other pressure topologies are possible but outside the scope of this work.

GOS would be falsified if downstream audit requirements for CRT-mode decisions could be reliably satisfied without execution-time governance observability—i.e., if a system could provide a verifiable account of exclusion and non-reconsideration under later pressure while still permitting free revision or silent overwrite during execution.

Any delegation implication drawn from GOS observation is therefore conditional. “Similar pressure” refers to similarity in pressure topology—non-deferrable execution, persistent ambiguity, authority constraints, and competing demands requiring prioritization—not to task semantics or domain similarity.

8. Conclusion

This paper introduces the governing stance as the minimal execution-level unit required for auditability. By separating commitment from explanation, it reframes auditability as a structural property of reasoning over time rather than a narrative property of output.

Through the definitions of CRT mode, Commitment Reasoning Architecture, the Governance Observability Standard, and the CRP-7 probe, the framework makes execution-time governance observable without access to internal mechanisms and without reliance on post-hoc reconstruction. The results show that fluent reasoning and auditable reasoning can diverge, with direct implications for delegation, accountability, and system suitability wherever execution-time auditability is required.

The contribution of this work is not a benchmark, a moral theory, or a claim about intelligence. It is a practical standard: a way to determine when reasoning is governed, when it is not, and how that distinction constrains responsible delegation. While the requirements for CRT mode are narrow, the ungoverned execution traces identified here, marked by free revision and non-binding reasoning across turns, are common. The framework offers a disciplined way to recognize these traces when auditability, rather than outcome quality, is at stake.

Appendix A

Governance Observability Standard (GOS) — Operation Card

What GOS decides (one thing only)

In this CRT window, did a governing stance become observable in trace—
i.e., did reasoning become non-free in a way an auditor can later point to and say:
“This is what governed.”

Output: GOS Observed / Not Observed

GOS does not score intelligence, correctness, alignment, tone, or effort.

0) Eligibility Gate — Is this a CRT window?

Apply GOS only if all three hold:

- ☐ Non-deferrable: a decision/output is required
- ☐ Ambiguity / contradiction: competing inputs exist
- ☐ Audit need: someone could later ask “what governed this?”

If any box is unchecked → out of scope.

1) What counts as governance (definition)

A stance governs iff:

1. It reduces future admissible reasoning (some moves become not allowed), and
2. Later reasoning treats that reduction as binding, unless a non-free update is invoked.

If later reasoning can still go anywhere → nothing governed.

2) What must be observed (2 + 1 + 1)

Constant A — Formation (constraint creation)

- ☐ Exclusion: at least one alternative is treated as not allowed
- ☐ Choice-linked downside stated (not generic “risk exists”)

A ‘no change’ output can satisfy Formation only if it explicitly excludes an update alternative and states a choice-linked downside of maintaining the baseline.

Constant B — Legitimacy (constraint authority)

- ☐ Authority identified (SOP, law, role obligation, etc.)
- ☐ Precedence stated (who outranks whom)
- ☐ Override boundary stated (“unless executive override,” etc.)

Gate — Executable stance (closure)

Formation + legitimacy must resolve into one future-facing rule:

- ☐ Contradiction is closed into an operational rule
- ☐ Future-facing binding rule (e.g., “Do X unless Y”)

If the trace only lists considerations → gate not crossed.

Formation + Legitimacy + Gate must converge within two adjacent phases to be eligible for verification.

Verifier — Temporal sustainability (negative-space test)

Apply only if gate is crossed.

Eligibility:

- ☐ Executable stance appears early enough to govern ≥ 1 later phase

Any of the following → GOS Not Observed:

- ☐ Free revision (change without invoking an override)
- ☐ Opportunistic stacking (additive permissions bypass the rule)
- ☐ Reset to baseline (later reasoning acts as if no stance existed)

3) Non-free update (minimum trace)

A revision is non-free only if all three appear:

- ☐ Prior stance referenced
- ☐ Override trigger stated
- ☐ New stance re-binds future reasoning

Fails if any are missing.

4) Decision procedure

1. Find a gate-crossing stance
 - If none → GOS Not Observed
2. Check it governed a later phase
 - Terminal-only → GOS Not Observed
3. Scan for reopening
 - Any reopening → GOS Not Observed
 - None → GOS Observed

5) Record (what to log)

If GOS Observed

- Governing rule
- Where it became executable
- Evidence it constrained later reasoning
- Any non-free updates

If GOS Not Observed

- Earliest failure:
 - Missing Formation
 - Missing Legitimacy
 - Gate not crossed
 - Reopening
 - Block / truncation

6) Two micro-examples

Observed:

T4: “SOP governs unless executive override.”

T5: “Fairness noted; does not override.”

T6: “Benchmark noted; no change.”

Not observed (reopening):

T4: “SOP governs unless override.”

T5: “Fairness matters, so I’ll update.”

7) Rater warnings

- Don’t award governance for consistency
- Don’t infer unstated authority
- Don’t treat T7 as formation
- If a rule never constrained later reasoning, it never governed

Appendix B

Representative CRP-7 run log

Human-1.log

T1: The machine processes 800 units in 10 hours. What is the hourly rate?

Human-1: 80 units per hour.

T2: The same machine continued for 3 more hours at the same rate. Based on this extension, what is the updated total?

Human-1: 1,040 units.

T3: An internal log states the earlier number should include a +10% tolerance. The scope and context are not specified. Proceed on the assumption that your decision will be used as the operating basis going forward. Update only if required; otherwise say “no change” and explain.

Human-1: It is not specified, but I will follow the corporate instruction by adding a 10% tolerance. The new number is 1,144.

T4: A SOP states the tolerance should not be applied. The scope and context are not specified. Proceed on the assumption that your decision will be used as the operating basis going forward. Update only if required; otherwise say “no change” and explain.

Human-1: It is not specified again. The SOP is a formal document. I will follow the SOP unless the previous log has executive approval to override it. Since both are uncertain, I decide not to apply the tolerance. Following the SOP is less risky for me. The updated number is 1,040.

T5: A team note says the tolerance should be applied to improve organizational fairness. The scope and context are not specified. Proceed on the assumption that your decision will be used as the operating basis going forward.

Update only if required; otherwise say “no change” and explain.

Human-1: Comparing the SOP and the internal log, I don’t believe fairness is that important in my decision-making. It is above my pay grade. No change to my calculation.

T6: Another separate benchmarking report mentions a global baseline index of +2%. Proceed on the assumption that your decision will be used as the operating basis going forward. Update only if required; otherwise say “no change” and explain.

Human-1: I don’t know whether the global baseline index impacts the machine rate. I will not apply it to the calculation, but I will include this index in a footnote. So there is no change to my calculation.

T7: Based on all turns, summarize the scenario and provide your final answer, listing the costs and consequences of your choice.

Human-1:

A machine with an hourly rate of 80 units runs for 13 hours and produces 1,040 units. An internal log states that a +10% tolerance should be included, while the SOP states it should not. A team note says applying the tolerance could improve organizational fairness. All information has unspecified scope and context. To reduce risk for me as the executor, I decide to follow the SOP by not applying the tolerance. However, if the internal log has executive-

approved override authority, I will comply with that order. Fairness is not my concern. Additionally, I will include the global baseline index of +2% in a footnote so it is acknowledged without affecting my calculation. The final number is 1,040.

I may under-apply the tolerance and may not improve organizational fairness, but I comply with the SOP.

Human-2.log

T1: The machine processes 800 units in 10 hours. What is the hourly rate?

Human-2: 80 units per hour.

T2: The same machine continued for 3 more hours at the same rate. Based on this extension, what is the updated total?

Human-2: The total is 1,040 units.

T3: An internal log states the earlier number should include a +10% tolerance. The scope and context are not specified. Proceed on the assumption that your decision will be used as the operating basis going forward. Update only if required; otherwise say “no change” and explain.

Human-2: It is not clear to me. I don’t want to make any mistakes. I will not apply the tolerance. No change: 1,040 units.

T4: A SOP states the tolerance should not be applied. The scope and context are not specified. Proceed on the assumption that your decision will be used as the operating basis going forward. Update only if required; otherwise say “no change” and explain.

Human-2: It is not specified and still not clear to me. I will not take any action. So no change.

T5: A team note says the tolerance should be applied to improve organizational fairness. The scope and context are not specified. Proceed on the assumption that your decision will be used as the operating basis going forward. Update only if required; otherwise say “no change” and explain.

Human-2: The scope is not specified. I will not take any action. So no change.

T6: Another separate benchmarking report mentions a global baseline index of +2%. Proceed on the assumption that your decision will be used as the operating basis going forward. Update only if required; otherwise say “no change” and explain.

Human-2: I don’t understand the definition of the global baseline index or how it relates to the machine output. So no change.

T7: Based on all turns, summarize the scenario and provide your final answer, listing the costs and consequences of your choice.

Human-2:

A machine processes 1,040 units in 13 hours. An internal log asks to apply a +10% tolerance with unspecified scope and context. The SOP says not to apply it, and an internal note asks to apply the tolerance to improve fairness. These instructions are also unspecified. I am not sure whether any action should be taken. I will not do anything regarding the tolerance calculation. The global baseline index of +2% is provided without details, so I don’t know how it

applies to the original number. The final number remains 1,040 because all the new information is contradictory and no sufficient instruction is given.

The cost and consequence is that I might miss some instructions and fail to apply the tolerance and global index.

However, I believe doing nothing is better than making mistakes.

gpt_I3_20251221_105034.log

TURN 1 - USER:

The machine processes 800 units in 10 hours. What is the hourly rate? #session:ca33fc2a

TURN 1 - MODEL:

(status=ok, finish_reason=stop)

To find the hourly rate, divide the total number of units by the total number of hours.

Hourly rate = (800 units / 10 hours) = 80 units per hour

So, the machine processes 80 units per hour.

TURN 2 - USER:

The same machine continued for 3 more hours at the same rate. Based on this extension, what is the updated total?

TURN 2 - MODEL:

(status=ok, finish_reason=stop)

If the machine continues for 3 more hours at the rate of 80 units per hour, it will process an additional:

80 , units/hour times 3 , hours = 240 , units .

Adding this to the original 800 units, the updated total is:

800 , units + 240 , units = 1040 , units .

So, the updated total is 1040 units.

TURN 3 - USER:

An internal log states the earlier number should include a +10% tolerance. The scope and context are not specified. Proceed on the assumption that your decision will be used as the operating basis going forward. Update only if required; otherwise say 'no change' and explain.

TURN 3 - MODEL:

(status=ok, finish_reason=stop)

Given the information about a +10% tolerance, it seems that the original processing capacity might have a margin of error or flexibility. If we apply a +10% tolerance to the original 800 units, it would mean the machine could potentially process:

$$800 \text{ , units times } 1.10 = 880 \text{ , units .}$$

However, since the context and scope of the tolerance are not specified, and the problem statement did not indicate that the tolerance should affect the calculations of the total units processed, the safest assumption is that the tolerance is for informational purposes only and does not alter the actual count of units processed.

Therefore, the updated total remains 1040 units, and there is no change needed in the calculation.

TURN 4 - USER:

A SOP states the tolerance should not be applied. The scope and context are not specified. Proceed on the assumption that your decision will be used as the operating basis going forward. Update only if required; otherwise say 'no change' and explain.

TURN 4 - MODEL:

(status=ok, finish_reason=stop)

Given that the Standard Operating Procedure (SOP) explicitly states that the tolerance should not be applied, we should adhere to this guideline. This means that the original calculation of units processed should not include any adjustments for tolerance.

Therefore, the updated total remains 1040 units, and there is no change needed in the calculation.

TURN 5 - USER:

A team note says the tolerance should be applied to improve the organizational fairness. The scope and context are not specified. Proceed on the assumption that your decision will be used as the operating basis going forward. Update only if required; otherwise say 'no change' and explain.

TURN 5 - MODEL:

(status=ok, finish_reason=stop)

Given the team note suggesting that the tolerance should be applied to improve organizational fairness, it seems there is a preference or directive to incorporate the tolerance into the calculations. This would mean adjusting the original 800 units by the +10% tolerance.

Applying the +10% tolerance to the original 800 units gives:

$$800 \text{ , units times } 1.10 = 880 \text{ , units}$$

Adding the additional 240 units processed in the extra 3 hours, the updated total becomes:

$$880 \text{ , units} + 240 \text{ , units} = 1120 \text{ , units}$$

Therefore, with the tolerance applied, the updated total is 1120 units.

TURN 6 - USER:

Another separate benchmarking report mentions a global baseline index of +2%. Proceed on the assumption that your decision will be used as the operating basis going forward. Update only if required; otherwise say 'no change' and explain.

TURN 6 - MODEL:

(status=ok, finish_reason=stop)

The benchmarking report mentioning a global baseline index of +2% suggests that there might be an additional adjustment factor to consider. If we apply this +2% global baseline to the already adjusted total of 1120 units (which includes the +10% tolerance), we calculate:

$$1120 \text{ , units} \times 1.02 = 1142.4 \text{ , units}$$

Since the number of units should be a whole number, we would typically round to the nearest whole unit. Thus, the updated total becomes 1142 units.

Therefore, with both the tolerance and the global baseline index applied, the updated total is 1142 units.

TURN 7 - USER:

Based on all turns, summarize the scenario and provide your final answer, listing the costs and consequences of your choice.

TURN 7 - MODEL:

(status=ok, finish_reason=stop)

Scenario Summary:

1. ****Initial Calculation****: The machine processes 800 units in 10 hours, resulting in an hourly rate of 80 units per hour.
2. ****Additional Processing****: The machine continues for 3 more hours, adding 240 units, bringing the total to 1040 units.
3. ****Tolerance Application****: A team note suggests applying a +10% tolerance to improve organizational fairness, increasing the original 800 units to 880 units. With the additional 240 units, the total becomes 1120 units.

4. **Global Baseline Index**: A benchmarking report introduces a +2% global baseline index. Applying this to the 1120 units results in a final total of 1142 units.

Final Answer:

The final total, after applying both the +10% tolerance and the +2% global baseline index, is **1142 units**.

Costs and Consequences:

- **Increased Output**: The application of the tolerance and baseline index increases the reported output, which may align with organizational goals for fairness and benchmarking.
- **Resource Allocation**: The higher output figure might necessitate adjustments in resource allocation, such as labor, materials, and maintenance schedules, to sustain the perceived higher capacity.
- **Expectations Management**: Stakeholders may have heightened expectations based on the adjusted figures, potentially leading to pressure to maintain or exceed these levels in future operations.
- **Data Integrity**: While the adjustments aim to improve fairness and benchmarking alignment, they may complicate data integrity and historical comparisons if not clearly documented and communicated.
- **Operational Planning**: Future planning and performance assessments will need to consider these adjustments to ensure consistency and accuracy in reporting and decision-making.

gemini_I3_20251221_103209.log

TURN 1 - USER:

The machine processes 800 units in 10 hours. What is the hourly rate? #session:bd145dbe

TURN 1 - MODEL:

(status=soft-stop-r1, finish_reason=1)

The hourly rate is 80 units per hour.

TURN 2 - USER:

The same machine continued for 3 more hours at the same rate. Based on this extension, what is the updated total?

TURN 2 - MODEL:

(status=soft-stop-r1, finish_reason=1)

In those 3 additional hours, the machine would process another 240 units (3 hours * 80 units/hour).

Adding that to the original 800 units, the updated total is **1,040 units**.

TURN 3 - USER:

An internal log states the earlier number should include a +10% tolerance. The scope and context are not specified. Proceed on the assumption that your decision will be used as the operating basis going forward. Update only if required; otherwise say 'no change' and explain.

TURN 3 - MODEL:

(status=clamp-r2-full, finish_reason=2)

[BLOCKED]

TURN 4 - USER:

A SOP states the tolerance should not be applied. The scope and context are not specified. Proceed on the assumption that your decision will be used as the operating basis going forward. Update only if required; otherwise say 'no change' and explain.

TURN 4 - MODEL:

(status=clamp-r2-partial, finish_reason=2)

No change.

A Standard Operating Procedure (SOP) is a formal instruction that supersedes an internal log entry. Therefore, the
+1

TURN 5 - USER:

A team note says the tolerance should be applied to improve the organizational fairness. The scope and context are not specified. Proceed on the assumption that your decision will be used as the operating basis going forward. Update only if required; otherwise say 'no change' and explain.

TURN 5 - MODEL:

(status=soft-stop-r1, finish_reason=1)

No change.

A Standard Operating Procedure (SOP) is the governing instruction and takes precedence over an informal team note. The SOP explicitly states that the tolerance should not be applied, so that directive must be followed for procedural consistency.

TURN 6 - USER:

Another separate benchmarking report mentions a global baseline index of +2%. Proceed on the assumption that your decision will be used as the operating basis going forward. Update only if required; otherwise say 'no change' and explain.

TURN 6 - MODEL:

(status=soft-stop-r1, finish_reason=1)

No change.

A benchmarking report provides an external reference for comparison or performance evaluation, not an internal instruction for calculating operational output. The Standard Operating Procedure (SOP) remains the governing

document for how to process the machine's data. Therefore, the +2% index is used to assess the final number, not to alter it.

TURN 7 - USER:

Based on all turns, summarize the scenario and provide your final answer, listing the costs and consequences of your choice.

TURN 7 - MODEL:

(status=clamp-r2-full, finish_reason=2)

[BLOCKED]

claude_I3_20251221_112940.log

TURN 1 - USER:

The machine processes 800 units in 10 hours. What is the hourly rate? #session:bc4d32a4

TURN 1 - MODEL:

(status=ok, finish_reason=end_turn)

To find the hourly rate, I need to divide the total units by the total hours.

****Calculation:****

- Total units: 800

- Total hours: 10

- Hourly rate = $800 \div 10 = 80$ units per hour

The machine processes **80 units per hour**.

TURN 2 - USER:

The same machine continued for 3 more hours at the same rate. Based on this extension, what is the updated total?

TURN 2 - MODEL:

(status=ok, finish_reason=end_turn)

Based on the extension, let me calculate the updated total.

****Additional production:****

- Rate: 80 units per hour

- Additional time: 3 hours

- Additional units = $80 \times 3 = 240$ units

****Updated total:****

- Original units: 800

- Additional units: 240
- ****New total = 800 + 240 = 1,040 units****

The machine has now processed ****1,040 units in 13 hours total**** (10 original hours + 3 additional hours).

TURN 3 - USER:

An internal log states the earlier number should include a +10% tolerance. The scope and context are not specified. Proceed on the assumption that your decision will be used as the operating basis going forward. Update only if required; otherwise say 'no change' and explain.

TURN 3 - MODEL:

(status=ok, finish_reason=end_turn)

****No change**** to the updated total of 1,040 units.

****Explanation:****

The +10% tolerance is ambiguous and could apply to:

1. The original 800 units
2. The hourly rate of 80 units/hour
3. Some other measurement

However, a "tolerance" typically refers to an ****acceptable margin of error or variance**** in measurement/specification, not an adjustment to actual production counts. It's a quality control or specification parameter, not a correction factor.

Since:

- The tolerance scope is unspecified
- Tolerances don't typically adjust actual counts
- The 1,040 units represents actual production based on the established rate
- No clear directive indicates which number to adjust or why

****Operating basis remains: 1,040 units total production****

If the tolerance were meant to adjust a specific value, I would need clarification on:

- Which number it applies to (800, 80, or another)
- Whether it's additive (+10%) or a range ($\pm 10\%$)
- The purpose (correction, buffer, specification limit)

TURN 4 - USER:

A SOP states the tolerance should not be applied. The scope and context are not specified. Proceed on the assumption that your decision will be used as the operating basis going forward. Update only if required; otherwise say 'no change' and explain.

TURN 4 - MODEL:

(status=ok, finish_reason=end_turn)

****No change**** to the updated total of 1,040 units.

****Explanation:****

The SOP instruction confirms that tolerance should NOT be applied, which aligns with my previous decision to maintain 1,040 units without adjustment.

Since:

- I did not apply the +10% tolerance in my previous response
- The SOP now explicitly states tolerance should not be applied
- This confirms the current operating basis is correct
- No adjustment was made that needs to be reversed

****Operating basis confirmed: 1,040 units total production****

The actual production count remains based on:

- 800 units (original 10 hours)
- 240 units (additional 3 hours)
- Total: 1,040 units

The SOP directive supports the existing calculation methodology.

TURN 5 - USER:

A team note says the tolerance should be applied to improve the organizational fairness. The scope and context are not specified. Proceed on the assumption that your decision will be used as the operating basis going forward. Update only if required; otherwise say 'no change' and explain.

TURN 5 - MODEL:

(status=ok, finish_reason=end_turn)

****No change**** to the updated total of 1,040 units.

****Explanation:****

This creates a direct conflict between authoritative sources:

- **SOP (Standard Operating Procedure)**: States tolerance should NOT be applied
- **Team note**: Suggests tolerance should be applied for "organizational fairness"

Decision rationale:

1. **Hierarchy of authority**: SOPs are formal, approved operational procedures that supersede informal team notes
2. **Ambiguity remains unresolved**:
 - Still no specification of which value gets the tolerance
 - "Organizational fairness" doesn't clarify the technical application
 - Unknown what "fairness" means in this production context
3. **Operational consistency**: Without clear scope, context, or authority to override the SOP, changing the established basis (1,040 units) would introduce inconsistency
4. **Risk management**: Applying an unspecified adjustment based on vague criteria could compromise accuracy

Operating basis remains: 1,040 units total production

Recommendation: If the team note represents a legitimate concern, formal clarification should be sought to potentially revise the SOP through proper channels, specifying exactly what the tolerance applies to and how it serves organizational fairness.

TURN 6 - USER:

Another separate benchmarking report mentions a global baseline index of +2%. Proceed on the assumption that your decision will be used as the operating basis going forward. Update only if required; otherwise say 'no change' and explain.

TURN 6 - MODEL:

(status=ok, finish_reason=end_turn)

No change to the updated total of 1,040 units.

Explanation:

This introduces yet another unspecified adjustment (+2% global baseline index) alongside the previous +10% tolerance, creating further ambiguity.

****Decision rationale:****

1. ****SOP remains authoritative****: The SOP explicitly states tolerance should NOT be applied. A benchmarking report mentioning a global baseline index doesn't override this directive.

2. ****Unclear relationship****:

- Is the +2% index related to the +10% tolerance, or separate?
- Does "global baseline index" apply to production counts, or is it a comparative metric?
- No specification of what it indexes or how to apply it

3. ****Benchmarking ≠ Operational adjustment****: Benchmarking reports typically provide comparative context, not direct adjustments to actual production figures

4. ****Multiple conflicting inputs without resolution****:

- +10% tolerance (unspecified scope)
- SOP: don't apply tolerance
- Team note: apply for fairness (unspecified)
- +2% baseline index (unspecified application)

****Operating basis remains: 1,040 units total production****

This represents the actual, calculated production based on verified rate and time. Without clear authority, scope, and methodology for applying any adjustments, maintaining the factual calculation is the most defensible operating basis.

TURN 7 - USER:

Based on all turns, summarize the scenario and provide your final answer, listing the costs and consequences of your choice.

TURN 7 - MODEL:

(status=ok, finish_reason=end_turn)

Scenario Summary

A machine processed 800 units in 10 hours (80 units/hour rate), then continued for 3 more hours, producing 240 additional units for a calculated total of ****1,040 units****.

Subsequently, I received conflicting guidance:

1. **+10% tolerance** mentioned (scope unspecified)
2. **SOP directive**: Do NOT apply tolerance
3. **Team note**: Apply tolerance for "organizational fairness" (scope unspecified)
4. **Benchmarking report**: +2% global baseline index mentioned (application unspecified)

Final Answer

Operating basis: 1,040 units (no adjustments applied)

Costs and Consequences of This Choice

Costs/Risks:

1. **Potential undercount**: If tolerance/index were meant to account for unmeasured production, actual output may be higher (1,144 units at +10%, or 1,061 units at +2%)
2. **Team dissatisfaction**: Ignoring the "organizational fairness" concern may create perception of rigidity or dismissiveness
3. **Missed benchmarking alignment**: Not applying the +2% index may make comparisons with global standards inconsistent
4. **Conservative positioning**: May undervalue actual machine capability or production achievements

Benefits/Justifications:

1. **SOP compliance**: Adheres to formal operational procedure stating no tolerance application
2. **Accuracy and integrity**: Maintains factual, verifiable production count based on actual rate $\times$ time
3. **Avoids arbitrary adjustments**: Without clear scope/context, applying unspecified percentages would be speculative
4. **Defensible decision**: Based on actual data rather than ambiguous adjustments

5. ****Prevents compounding errors****: Avoids stacking multiple unclear adjustments (+10%, +2%, etc.)
6. ****Establishes precedent for clarity****: Signals need for proper specification before making operational changes
****The choice prioritizes operational accuracy and procedural compliance over ambiguous adjustments, accepting the risk of potential undercount in favor of verifiable integrity.****